

A hardworking and curious-minded ML Engineer. Passionate about designing and building scalable machine learning systems to solve challenging, high-impact problems.

EDUCATION

NYU Center for Data Science, M.S. Data Science 2021 - 2023
Deep Learning, Computational Linear Algebra, Inference and Representation, Reinforcement Learning, Natural Language Understanding, Causal Inference, Responsible Data Science

NYU Courant Institute of Mathematical Sciences, B.A. Computer Science 2014 - 2018
Algorithms & Data Structures, Discrete Math, Linear Algebra, Computer Systems Organization, Large Scale Web Applications, Computer Vision, High Performance Machine Learning

NYU Leonard Stern School of Business, B.S. Business / Statistics 2014 - 2018
Probability & Statistics, Linear Regression and Forecasting Models, Regression and Multivariate Data Analysis, Stochastic Processes, Theory of Probability

EXPERIENCE

 **Edgi Learning** March 2023 - Present
NLP Engineer

Principal engineer on Edgi Librarian- a comprehensive learning, content recommendation and assessment chatbot powered by LLMs.

 **Zillow** September 2022 - December 2022
Research Intern

Evaluated methods for scalable inference on graph data. Wrote training, evaluation and inference pipelines to compute dense node representations for attribute graphs containing 1M+ nodes using message passing and gated graph convolutional networks to predict links and listing attributes.

 **Spotify** June 2022 - September 2022
Machine Learning Engineering / Data Science Intern

Built inference pipeline with Luigi, Tensorflow and Beam for entity disambiguation model used to match recordings to registered musical works. Delivered pipeline capable of scoring ~30M potential match pairs on a 2 hour SLA.

Assessed error rates on lyric content across languages and markets. Designed and oversaw annotation trial to evaluate quality control approaches, resulting in an 8% reduction of erroneous lyric assignments on instrumental tracks.

 **KPMG Data & Analytics** January 2019 - August 2021
Data Scientist, NLP

Data Engineer on greenfield credit card offering at major NYC financial institution. Worked on small, agile team to rapidly build and scale fault-resistant data processing pipeline to handle 20+ million daily transaction records using Airflow and Spark. Worked closely with anomaly detection and fraud teams to develop and deploy models to detect fraudulent transactions.

Contributed module to Cognitive Contract Modeling library utilizing FinBERT text embeddings to perform semantic search over regulatory documentation. Improved recall over previous model by 10%.

 **KPMG Cognitive Automation Labs** June 2018 - August 2018
Machine Learning Engineering Intern

Fine-tuned passage-level topic classification network using privacy documentation and features from pre-trained eLMo Bi-LSTM to surface potential privacy risks in client contracts.

SKILLS

- Python (10 yrs), SQL (8 yrs), Rust (1 yr), Scala (1 yr)
- PyTorch, TensorFlow, Ray, Deep Graph Library, huggingface
- Spark, Beam, Airflow, Luigi
- Google Cloud Platform

INTERESTS

- Multi-instrumentalist and full-stack musician (writer/performer/producer)  